

# Latent Diffusion (Stable-Diffusion) Model



Zhiping (Patricia) Xiao

UCLA SCAI Lab Course Reading Group Winter 2023

## Introduction

References

Background

## Latent Diffusion

Diffusion Model for Images

Motivation

Architecture

Training

Applications

## Conclusion

More on Latent/Stable Diffusion

Similar Works

Social Impact

# Introduction

---



High-Resolution Image Synthesis with Latent Diffusion Models  
(CVPR'22)

- ▶ Operating on **latent space** of pre-trained auto-encoders, instead of on **pixel space**.
- ▶ Code: <https://github.com/CompVis/latent-diffusion> and <https://github.com/CompVis/stable-diffusion>

A YouTube video from **Lightning AI**:

<https://www.youtube.com/watch?v=AQrMWH8aC0Q>

Prerequisites:

- ▶ U-Net: Convolutional Networks for Biomedical Image Segmentation (MICCAI'15)

Related Works:

- ▶ GLIDE: Towards Photorealistic Image Generation and Editing with Text-Guided Diffusion Models (CVPR'21)
- ▶ **(CLIP)** Learning Transferable Visual Models From Natural Language Supervision (CVPR'21)
- ▶ **(DALLE)** Zero-Shot Text-to-Image Generation
- ▶ **(DALLE-2)** Hierarchical Text-Conditional Image Generation with CLIP Latents (CVPR'22)

# Quick Recap: Generative Models

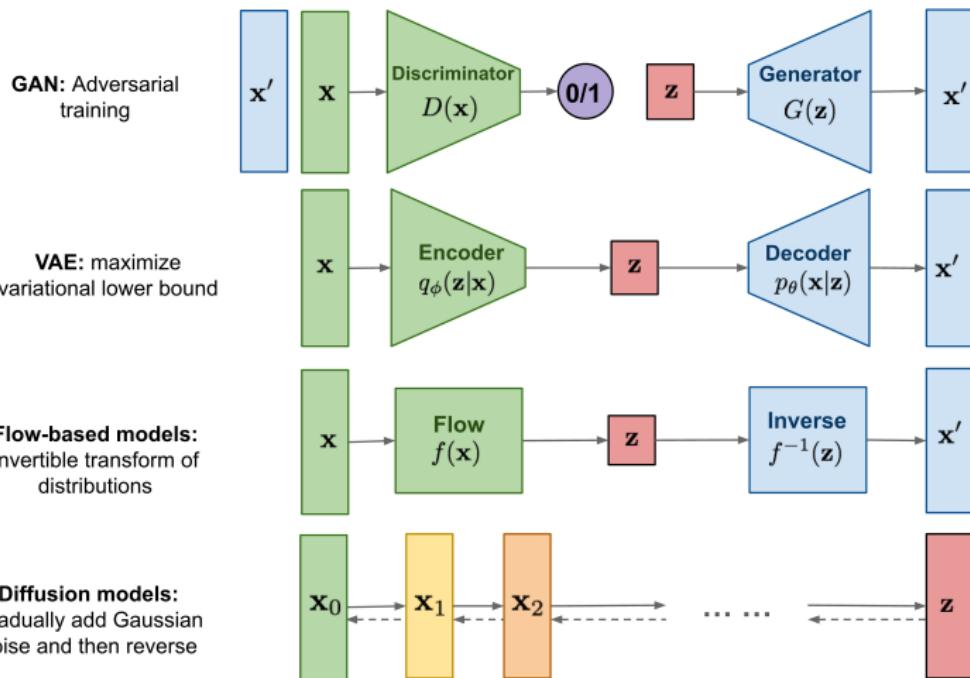


Figure: Overview of different types of generative models.<sup>1</sup>

<sup>1</sup>[lilianweng.github.io 2021 post on diffusion models](https://lilianweng.github.io/2021/post/diffusion-models/)

Table: Advantage &amp; Disadvantage of the Generative Models

Model	likelihood-based?	Good At	Not Good At
GAN	NO	efficient sampling; perceptual quality	optimize; capture data distribution
VAE/ Flow- based	YES	capture data distribution; opti- mize	perceptual qual- ity
DMs	YES	capture data dis- tribution; percep- tual quality	computation cost

To put it in other words:

- ▶ Capture Full Data Distribution: being able to generate unobserved samples, e.g. ridiculous photos
- ▶ Perceptual Quality: high fidelity, looks like real and much detailed, e.g. producing photo-realistic images
- ▶ Optimize: stable training, easy to optimize.

Earlier days: being able to identify objects in images.

2015: Automated Image Captioning (Images to Natural Language Descriptions)

Earlier days: being able to identify objects in images.

2015: Automated Image Captioning (Images to Natural Language Descriptions)

2016: Images from Captions

- ▶ “We can do image to text, why not try doing text to image and see how it works?” – Elman Mansimov (AWS)
- ▶ Generating Images from Captions with Attention (ICLR’16)
  - ▶ LM: Bidirectional Attention RNN
  - ▶ IM: Conditional DRAW Network (stochastic RNN that consists of a sequence of latent variables)
- ▶ Blurred but reasonable results. Being able to generate something unobserved before. e.g. Blue school bus.



A yellow school bus  
parked in a parking lot.



A red school bus parked  
in a parking lot.



A green school bus  
parked in a parking lot.



A blue school bus parked  
in a parking lot.



The decadent chocolate  
desert is on the table.



A bowl of bananas is on  
the table.



A vintage photo of a cat.



A vintage photo of a dog.

Figure 3: **Top:** Examples of changing the color while keeping the caption fixed. **Bottom:** Examples of changing the object while keeping the caption fixed. The shown images are the probabilities  $\sigma(c_T)$ . Best viewed in colour.

AI Art isn't new. e.g. Morphing portraits, style transfer, etc.

- ▶ Generating a specific type of image is easy. e.g. faces.
- ▶ Generating a scene from any natural language description?

AI Art isn't new. e.g. Morphing portraits, style transfer, etc.

- ▶ Generating a specific type of image is easy. e.g. faces.
- ▶ Generating a scene from any natural language description?

2021: OpenAI introduced DALL-E

- ▶ Usage: <https://github.com/openai/DALL-E>
- ▶ Source code NOT released.
- ▶ Not diffusion model yet. GPT-like model, auto-regressively generating an image from text + start of an image.

2021: GLIDE comes after: [arxiv.org/abs/2112.10741](https://arxiv.org/abs/2112.10741)

- ▶ Open-sourced  
<https://github.com/openai/glide-text2im>

AI Art isn't new. e.g. Morphing portraits, style transfer, etc.

- ▶ Generating a specific type of image is easy. e.g. faces.
- ▶ Generating a scene from any natural language description?

2021: OpenAI introduced DALL-E

- ▶ Usage: <https://github.com/openai/DALL-E>
- ▶ Source code NOT released.
- ▶ Not diffusion model yet. GPT-like model, auto-regressively generating an image from text + start of an image.

2021: GLIDE comes after: [arxiv.org/abs/2112.10741](https://arxiv.org/abs/2112.10741)

- ▶ Open-sourced  
<https://github.com/openai/glide-text2im>

Prompt Engineering: the craft of communicating with these zero-shot pre-trained deep learning models.

The first popular AI-painter that gets attention this year (2022), is DALL·E-2 (<https://openai.com/dall-e-2/>). With:

- ▶ Higher resolution;
- ▶ Greater comprehension;
- ▶ New capabilities e.g. in-painting.

And in a way, it actually helps reveal “how AI understands our world”. Triggered a trend of text-to-image generators, mostly diffusion models. (e.g. MidJourney)

The first popular AI-painter that gets attention this year (2022), is DALL·E-2 (<https://openai.com/dall-e-2/>). With:

- ▶ Higher resolution;
- ▶ Greater comprehension;
- ▶ New capabilities e.g. in-painting.

And in a way, it actually helps reveal “how AI understands our world”. Triggered a trend of text-to-image generators, mostly diffusion models. (e.g. MidJourney)

Latent (Stable) Diffusion is probably not the best model, but it has a great advantage being **open-sourced** (and released parameters). Therefore, everyone adapted their model and play with this version of text-to-image generator (e.g. <https://novelai.net/>).

- ▶ Note: can do much more than text-to-image generation.  
(Condition can be more than text.)

Time	Model	Availability
Dec 2021	Stable/Latent Diffusion	official code
Mar 2022	MidJourney	no paper yet
Apr 2022	DALL·E-2	unofficial code

Table: The star AI painting tools in year 2022.

# Latent Diffusion

---





Figure: Input: Random Noise (of the size of the image); Output Sequence: De-noised result from the previous step; Final Output: Image. *Trained via learning parameters to apply noise to images iteratively until it is complete noise, and inference by going through the opposite way.* (from video YouTube Video mentioning Imagen)

Diffusion Process: go from little noise to more noise.

Backward Diffusion Process: the reverse, from more noisy version to clearer image.

Why having  $T$  steps?

Diffusion Process: go from little noise to more noise.

Backward Diffusion Process: the reverse, from more noisy version to clearer image.

Why having  $T$  steps?

- ▶ Breaking-down a hard problem (generating images from noise).
- ▶ Easier to inject condition (e.g. text information) gradually than all at once.

Belong to the class of likelihood-based probabilistic models.

Learn a data distribution  $p(x)$  by gradually denoising a normally-distributed variable  $x_T$ , i.e. learning the **reverse process** of a fixed Markov Chain of length  $T$ .

$$L_{DM} = \mathbb{E}_{x, \epsilon \sim \mathcal{N}(0,1), t} \left[ \|\epsilon - \epsilon_\theta(x_t, t)\|_2^2 \right],$$

where  $t$  is uniformly sampled from  $\{1, 2, \dots, T\}$ , and the models can be interpreted as an equally-weighted sequence of denoising autoencoders  $\epsilon_\theta(x_t, t)$  ( $\epsilon_\theta$  is typically implemented as **U-Net**), trained to predict the denoised version of  $x_t$ , namely  $x$ . *The above version works well in **image** settings.*

Diffusion models are capable of modeling conditional distributions  $p(x|y)$  via using a conditional denoising autoencoder  $\epsilon_\theta(x_t, t, y)$ .

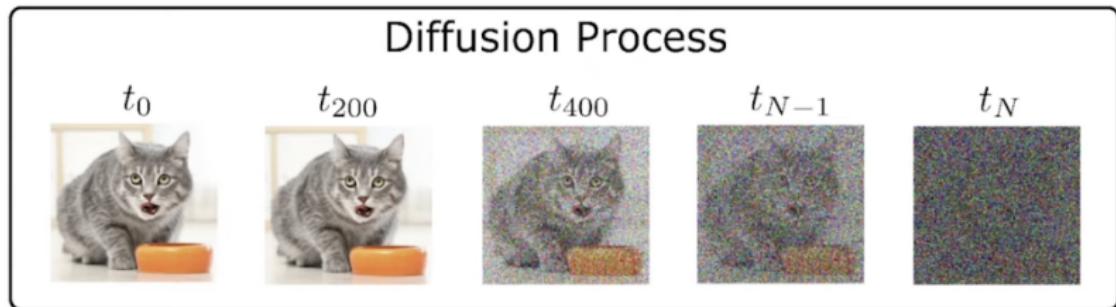


Figure: Diffusion Process. Illustration comes from Lightning AI YouTube video.

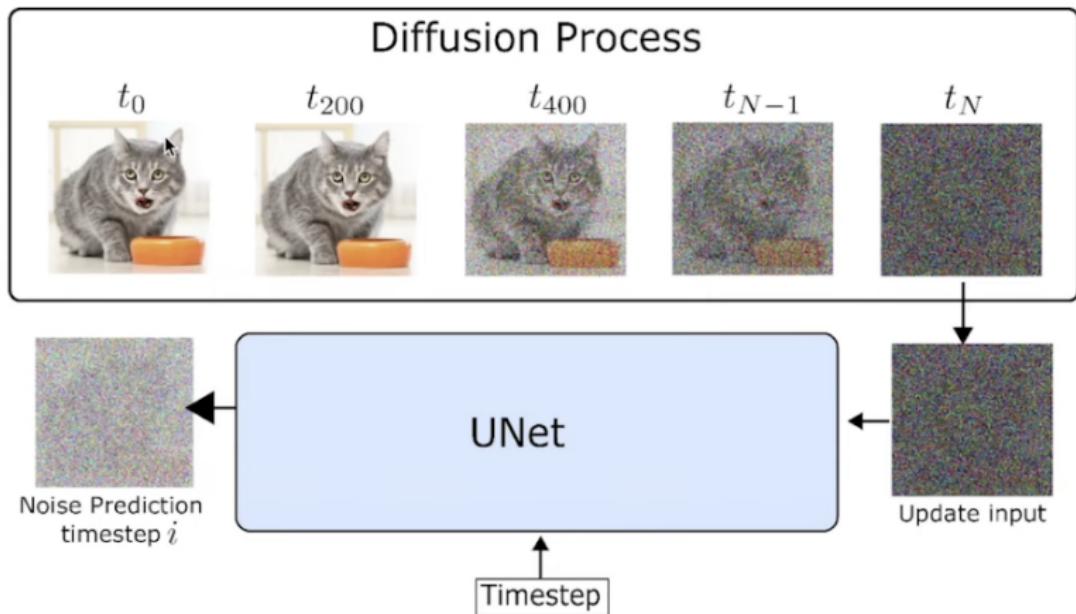


Figure: Backward Diffusion Process in general. Illustration comes from Lightning AI YouTube video.

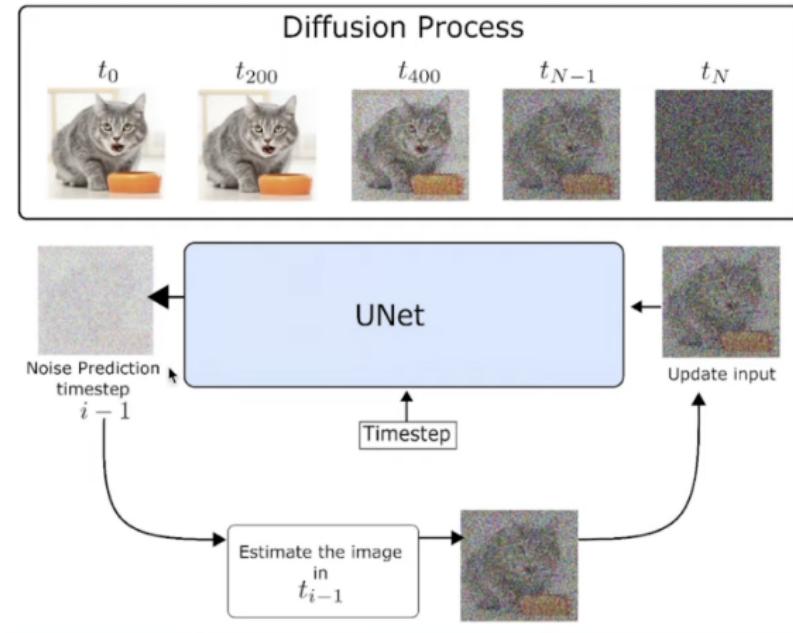


Figure: Backward Diffusion Process iteration  $i$ . Illustration comes from Lightning AI YouTube video.

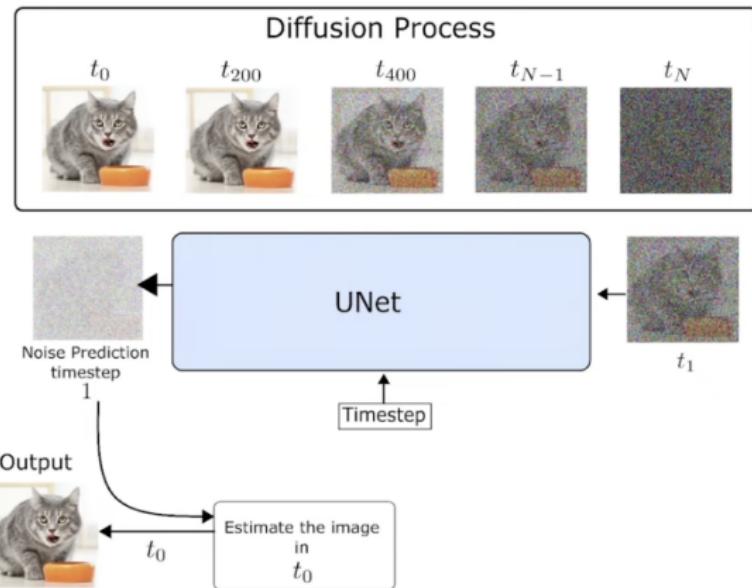


Figure: Backward Diffusion Process gets the final outcome.  
Illustration comes from Lightning AI YouTube video.

$$L_{DM} = \mathbb{E}_{x,\epsilon \sim \mathcal{N}(0,1),t} \left[ \|\epsilon - \epsilon_\theta(x_t, t)\|_2^2 \right]$$

The most powerful DMs are often computationally demanding.

- ▶ Costly Training: UNet has typically  $\approx 800M$  parameters; the model takes hundreds of GPU days to train, prone to spend excessive amounts of capacity on modeling imperceptible details;<sup>2</sup>
- ▶ Costly Evaluation: cost a lot of time and memory, must run the same architecture sequentially for many of steps.
- ▶ e.g. Diffusion Models Beat GANs on Image Synthesis (NeurIPS'21) takes 150 - 1000 *V100 days* to train, 25 - 1000 steps to evaluate.

<sup>2</sup>This is because of the Mode-Covering behavior (reference)

Minimize difference between  $Q(x)$  and  $P(x)$  at data distribution  
 $P(x) > 0$ :

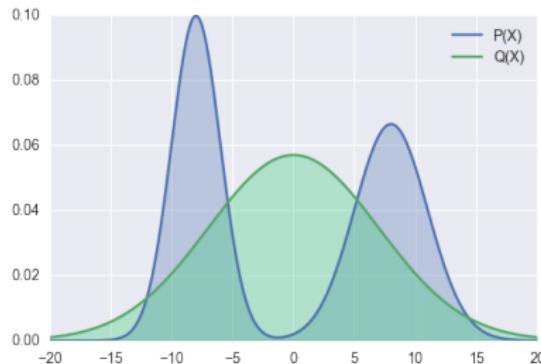
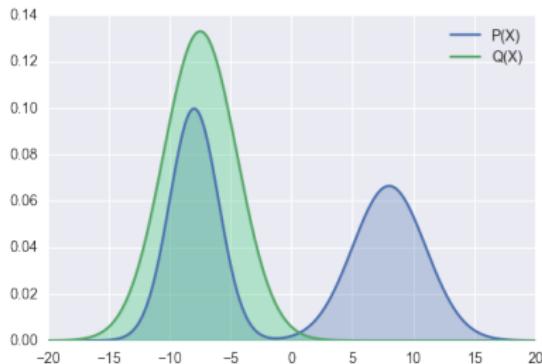


Figure: Bad v.s. Good mode-covering. On the left (bad) example, the right hand side **mode** is not covered by  $Q(x)$ , but it is the case that  $P(x) > 0$ . (from agustinus.kristia.de)

*This behavior makes the DMs costly.*

Previous solutions of the cost:

- ▶ Weighted importance of steps. (still expensive)

$$L_{DM} = \mathbb{E}_{x, \epsilon \sim \mathcal{N}(0,1), t} \left[ \lambda(t) \|\epsilon - \epsilon_\theta(x_t, t)\|_2^2 \right]$$

- ▶ Having an extra model to learn upsampling & sharpening of images. (still working on image space) e.g. GLIDE.

**Observation:** Most bits of a digital image correspond to imperceptible details, but we still need to train and evaluate on all pixels if we work on pixel spaces.

Highlighted Novelty: Do Diffusion on **Latent** Space, and accept more general types of conditions.

- ▶ Operating on latent space (perceptually equivalent space) of powerful pre-trained auto-encoders, instead of directly on pixel space (*compared with standard Diffusion Models*).
- ▶ **Less Costly:** Fast sampling, efficient training, one-step decoding to image space.
- ▶ **More Flexibility:** More general conditions. (Besides, operation on latent space makes it easier to add other signals.)

Three Major Components, trained **separately**:

- ▶ Autoencoder: Implemented as Variational Autoencoder (VAE); Handling perceptual image compression.
  - ▶ Encoder  $\mathcal{E}$ , Decoder  $\mathcal{D}$
  - ▶  $z = \mathcal{E}(x)$  where the RGB image  $x \in \mathbb{R}^{H \times W \times 3}$  turns into latent representation  $z \in \mathbb{R}^{h \times w \times c}$ , while  $\tilde{x} = \mathcal{D}(z)$  tries to reconstruct  $x$ .
  - ▶ Some regularization terms applied (e.g. KL-penalty towards a standard normal) to avoid arbitrarily high-variance.
- ▶ Denoiser: **Latent Diffusion Models**

$$L_{LDM} = \mathbb{E}_{\mathcal{E}(x), \epsilon \sim \mathcal{N}(0,1), t} \left[ \|\epsilon - \epsilon_\theta(z_t, t)\|_2^2 \right],$$

where the neural backbone  $\epsilon_\theta$  of LDM is realized as a time-conditional attention UNet.

- ▶ Conditioning Encoder: can be arbitrary encoder that produces a sequence of tokens.

DMs are capable of modeling conditional distribution  $p(z|y)$ , by using a conditional denoiser  $\epsilon_\theta(z_t, t, y)$ .

LDMs propose to augment the UNet backbone implementing  $\epsilon_\theta$  with the cross-attention mechanism.  $Q$ ,  $K$ ,  $V$  are projection of  $\phi_i(z_t)$ ,  $\tau_\theta(y)$ ,  $\tau_\theta(y)$  respectively, where  $\phi_i(z_t)$  is the flattened intermediate representation of the U-Net implementing  $\epsilon_\theta$ .

$$L_{LDM} = \mathbb{E}_{\mathcal{E}(x), \epsilon \sim \mathcal{N}(0,1), t} \left[ \|\epsilon - \epsilon_\theta(z_t, t, \tau_\theta(y))\|_2^2 \right],$$

$\tau_\theta$  is a domain specific encoder used to project  $y$ , e.g.  $\tau_\theta$  can be transformers when  $y$  are text prompts.

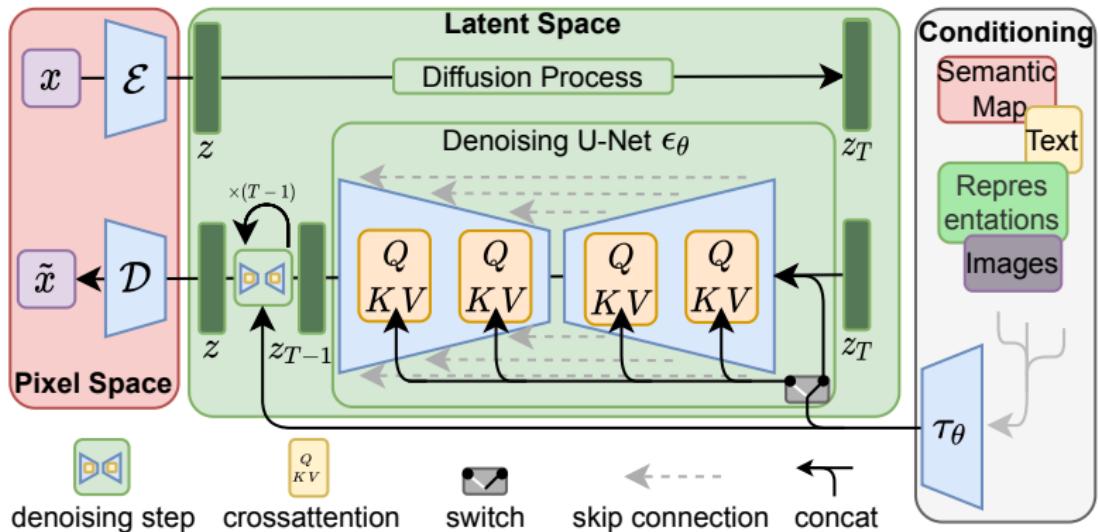


Figure: Figure 3 in the original paper.

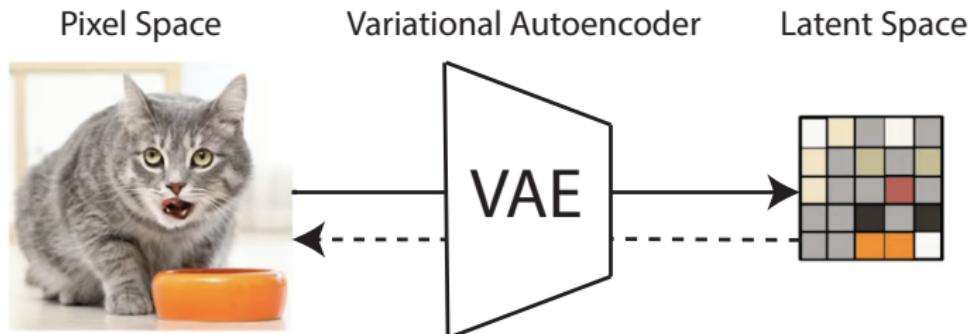
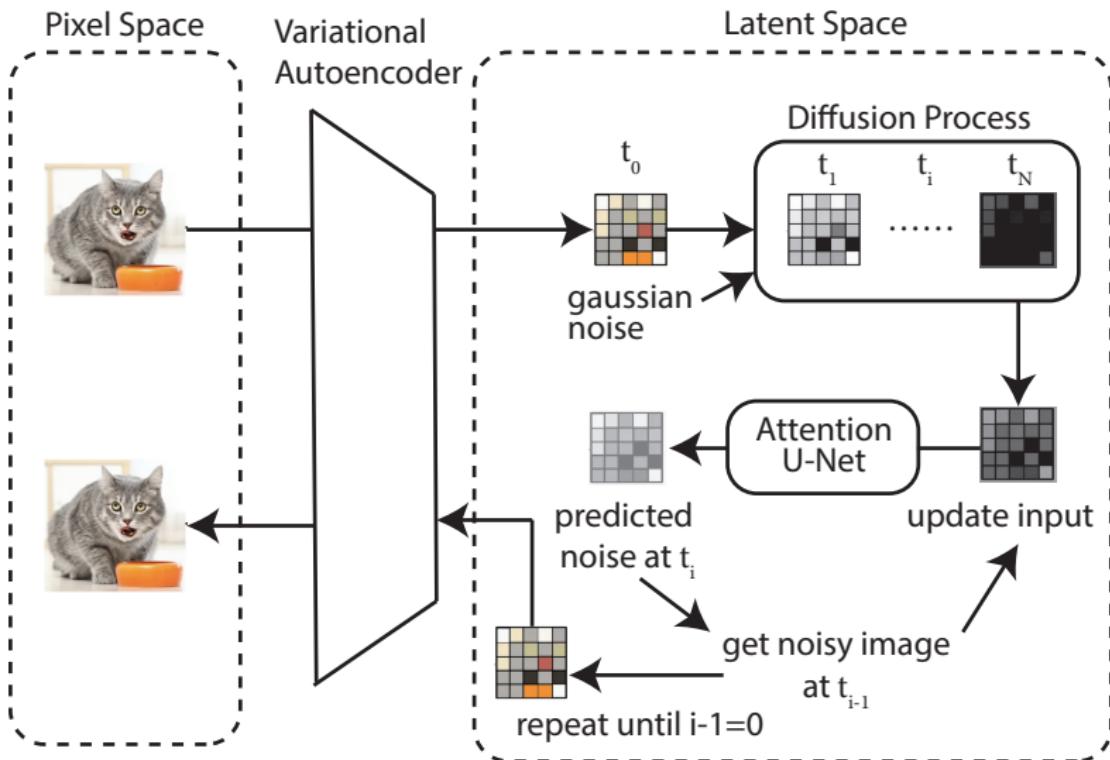
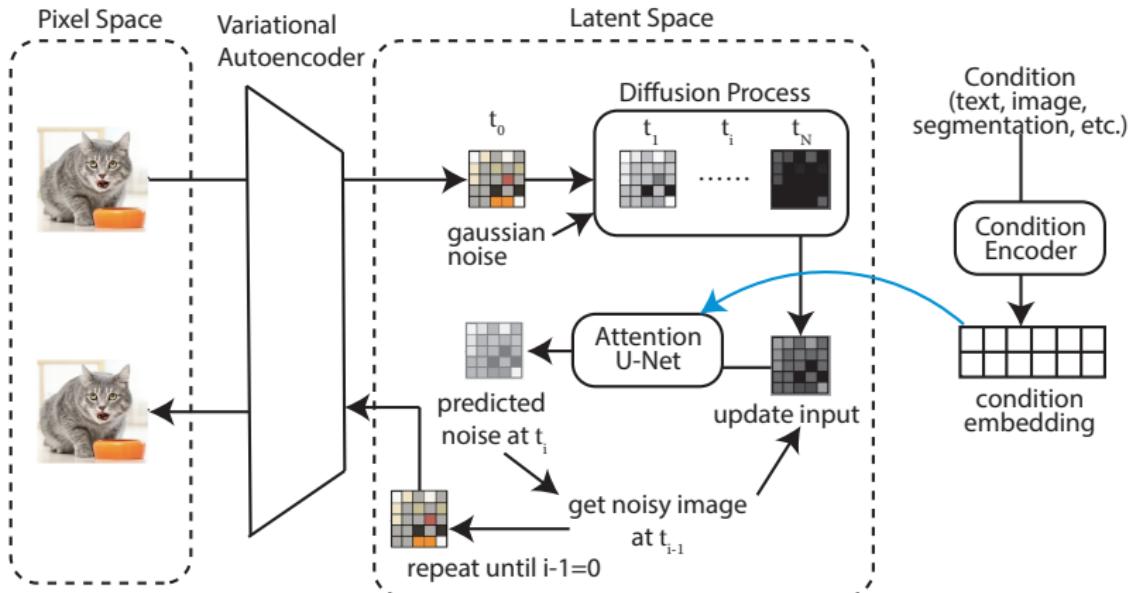


Figure: The role autoencoder plays in the model. Note: we actually use **separated** encoder  $\mathcal{E}$  and decoder  $\mathcal{D}$ . But this detail is not illustrated in this figure precisely for simplicity concern.





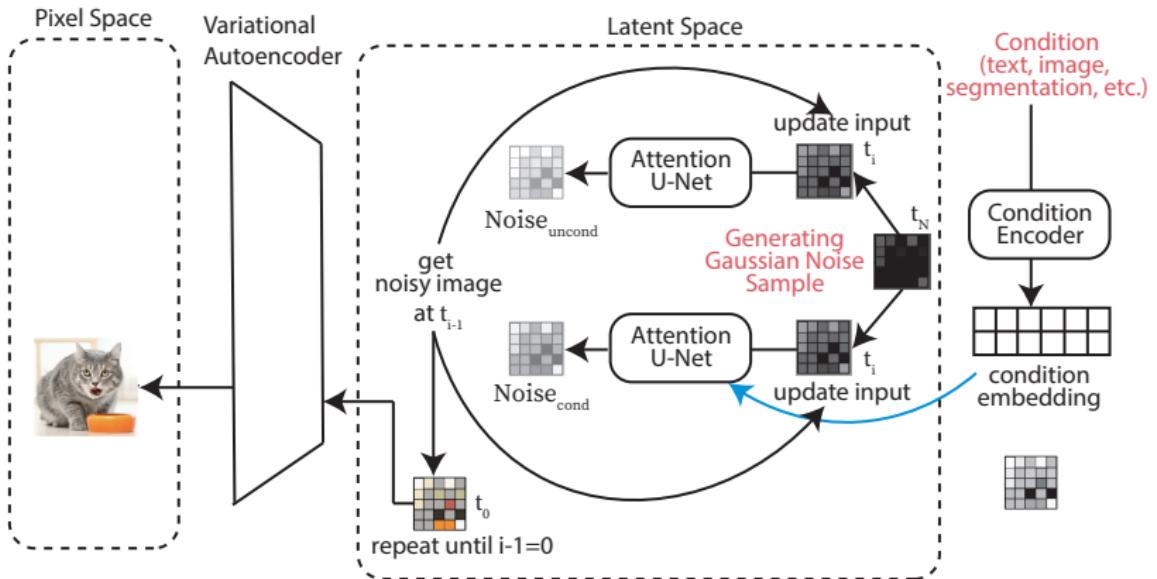


Figure:  $Noise = Noise_{uncond} + \beta(Noise_{cond} - Noise_{uncond})$

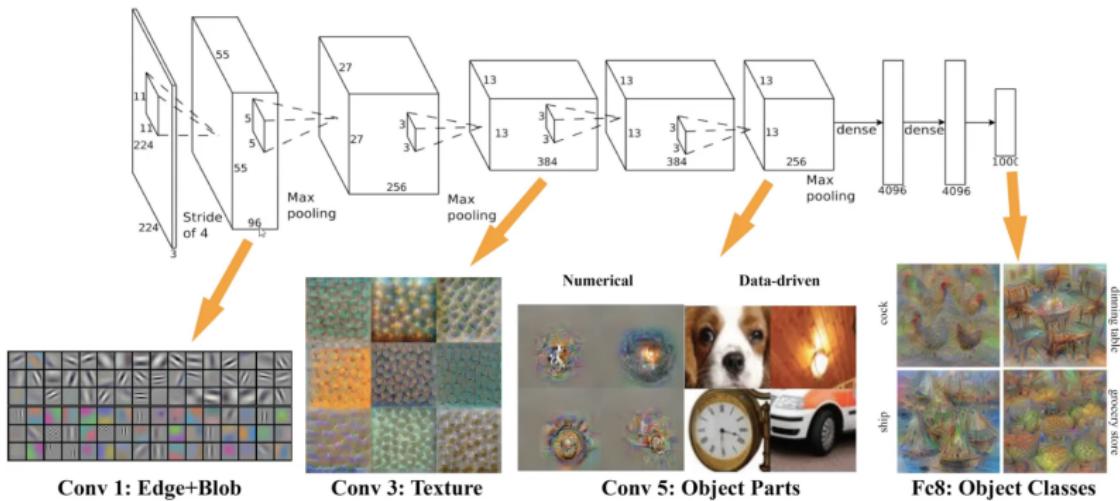


Figure: An example of a convolutional-based model. Image from YouTube. **U-Net** is a special kind of convolutional-based model.

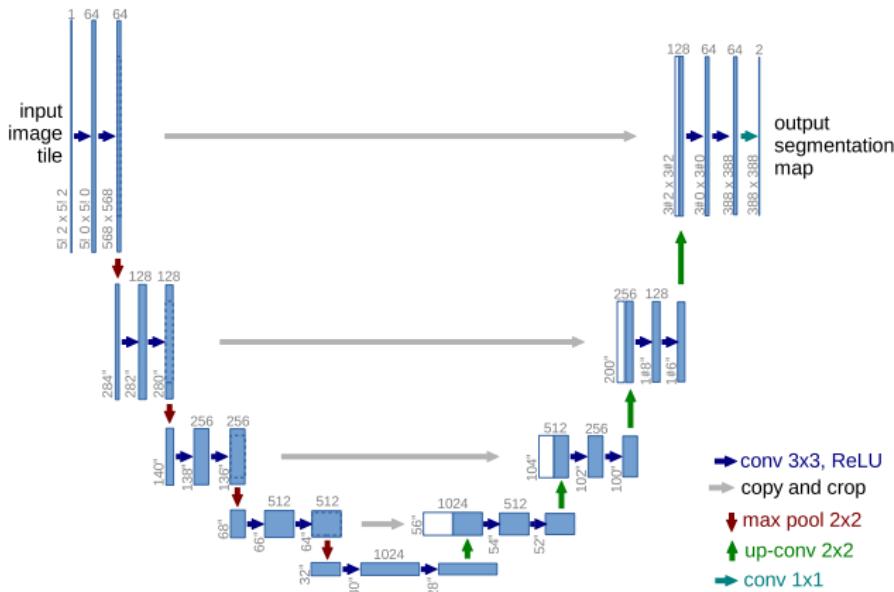


Figure: Blue box: a multi-channel feature map; White box: copied feature maps. Very common tool for image segmentation. Preserves the dimensionality.

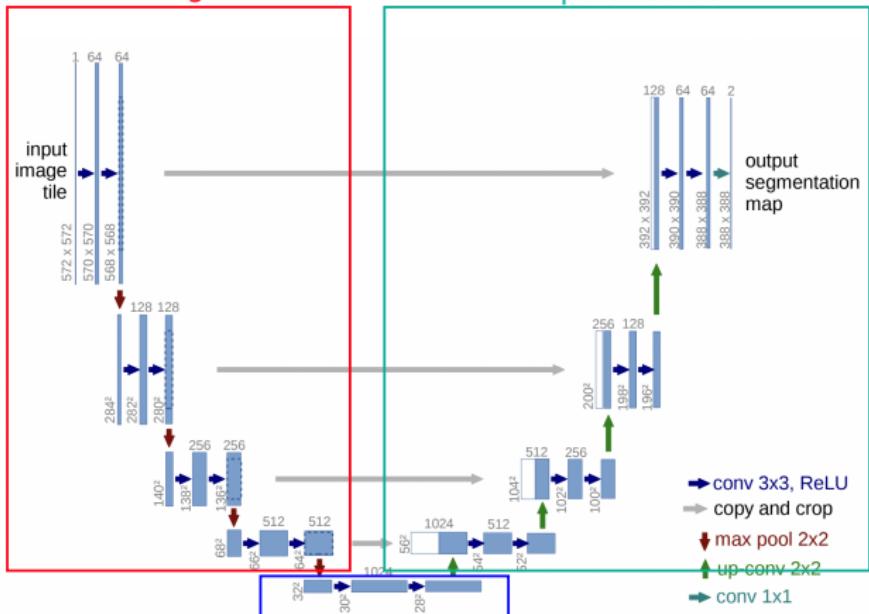
# Quick Recap: U-Net

35

decreasing size  
increasing feature  
capture context

increasing size  
decreasing feature  
combining contextual information  
enables localization  
**Expansive Block**

## Contracting Block



Bottleneck  
keep the same size  
increasing feature  
capture context

UCLA

## How LDM uses Unet?

LDM changes convolution + ReLu layers by **ResNet+Spatial Transformer** layers.

LDM adds timestep information and context information.

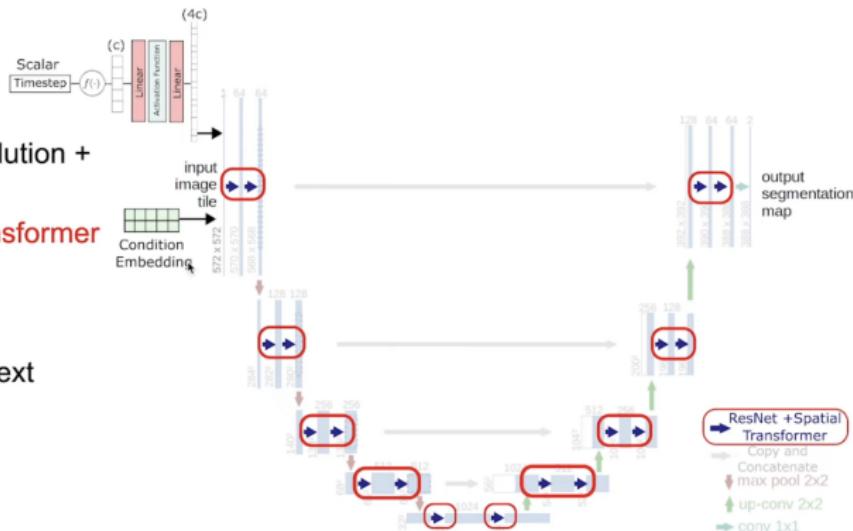


Figure: from Lightning AI video. *Context embedding* is the *condition embedding*.

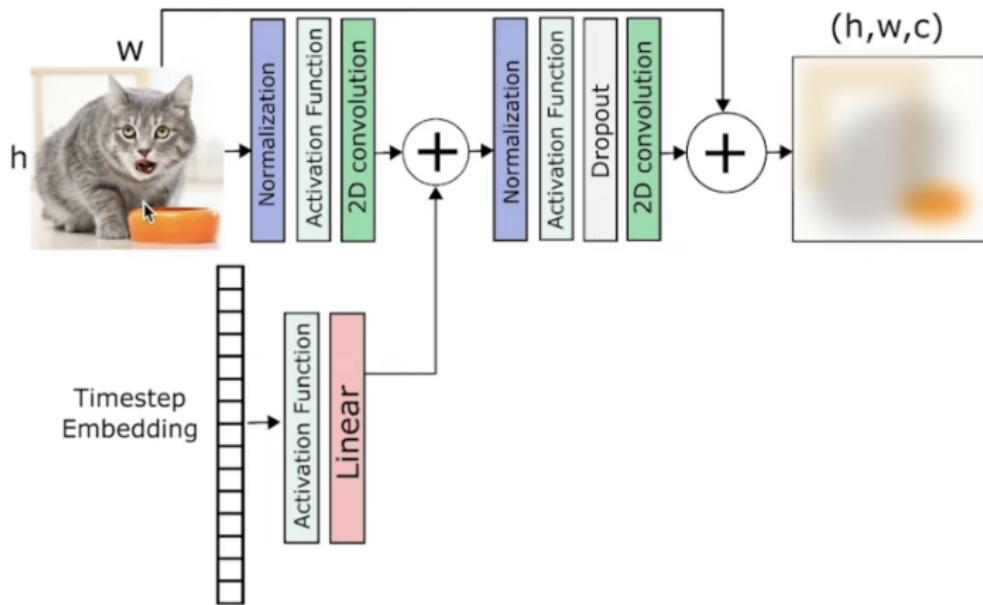


Figure: from Lightning AI video. Takes in time-step embedding.

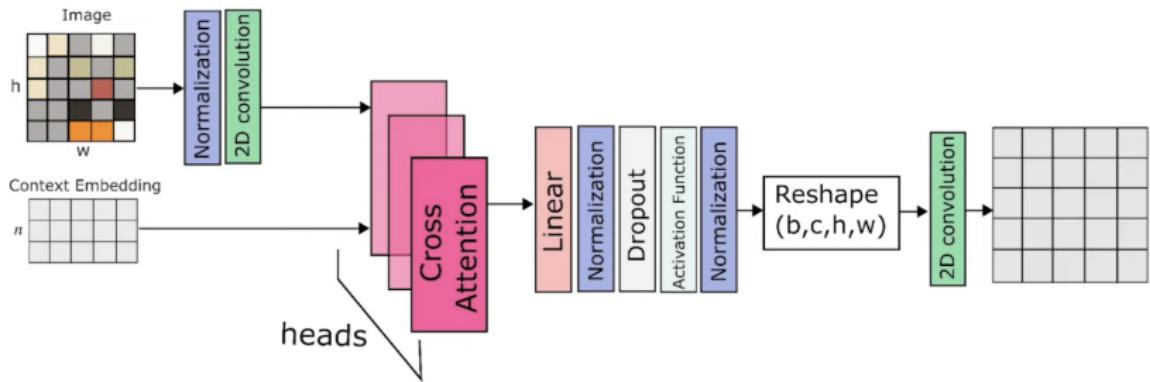


Figure: from Lightning AI video. Takes in context (condition) embedding.

# Quick Recap: Attention

39

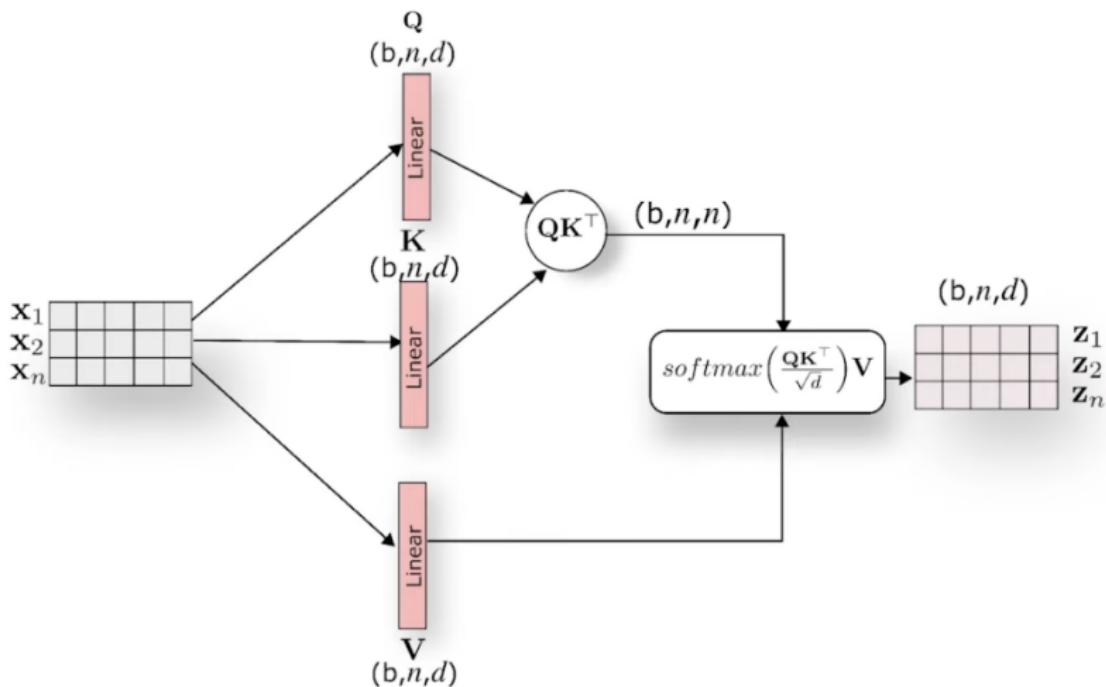


Figure: from Lightning AI video.

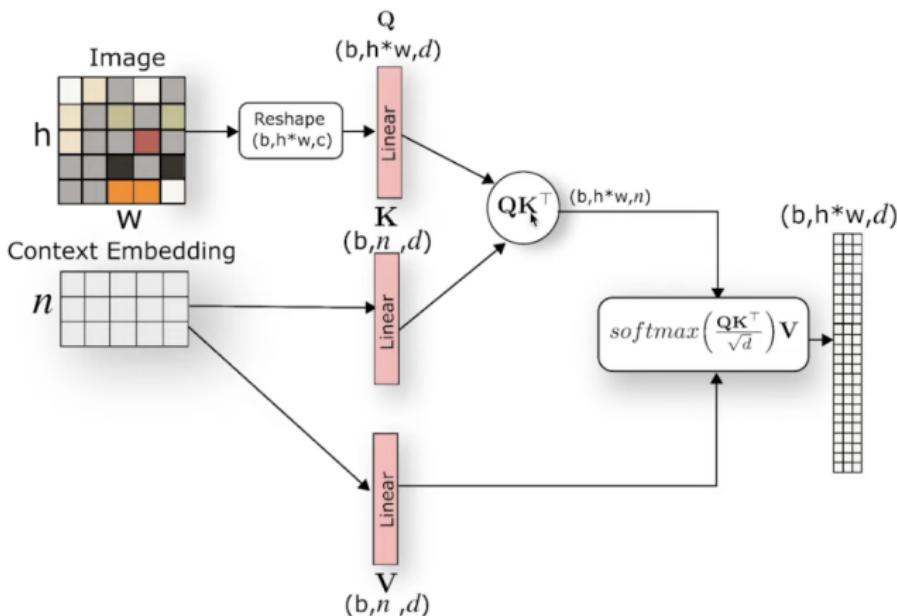
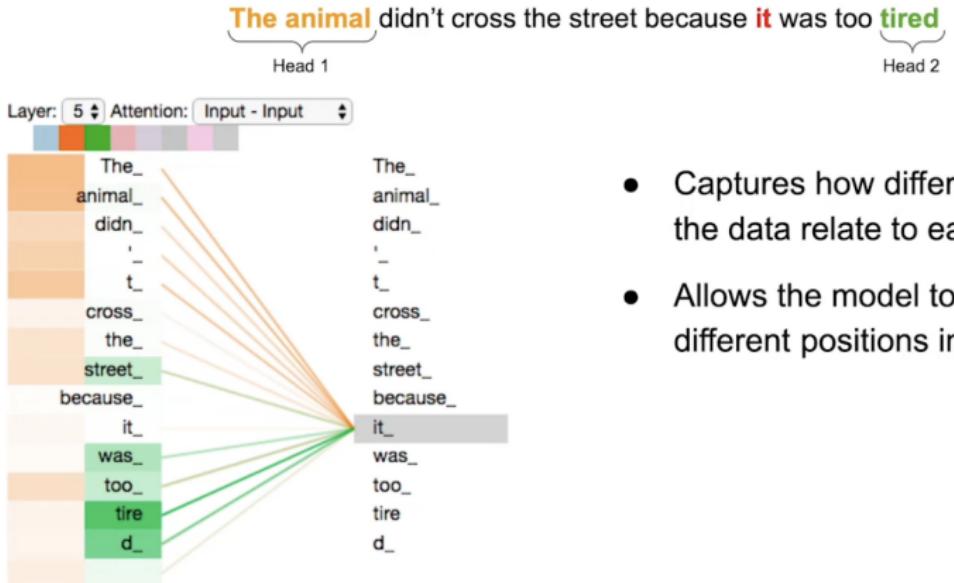


Figure: from Lightning AI video.  $Q$  from image embedding,  $K$  and  $V$  both from context (condition) embedding.

## Quick Recap: Multi-Head Attention



- Captures how different parts of the data relate to each other
  - Allows the model to focus in different positions in the data

**Figure:** from Lightning AI video. For more examples like this visit this blog.

A Hedgehog crossing the street



Figure: from Lightning AI video.

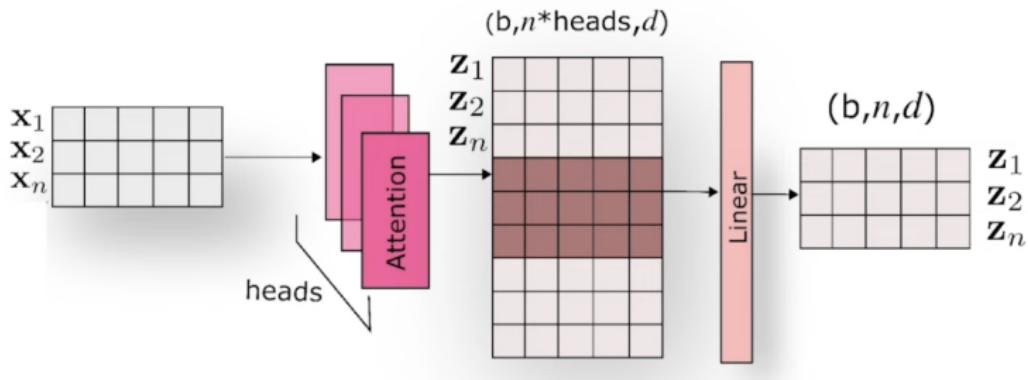


Figure: from Lightning AI video. Allows for more flexible attention.

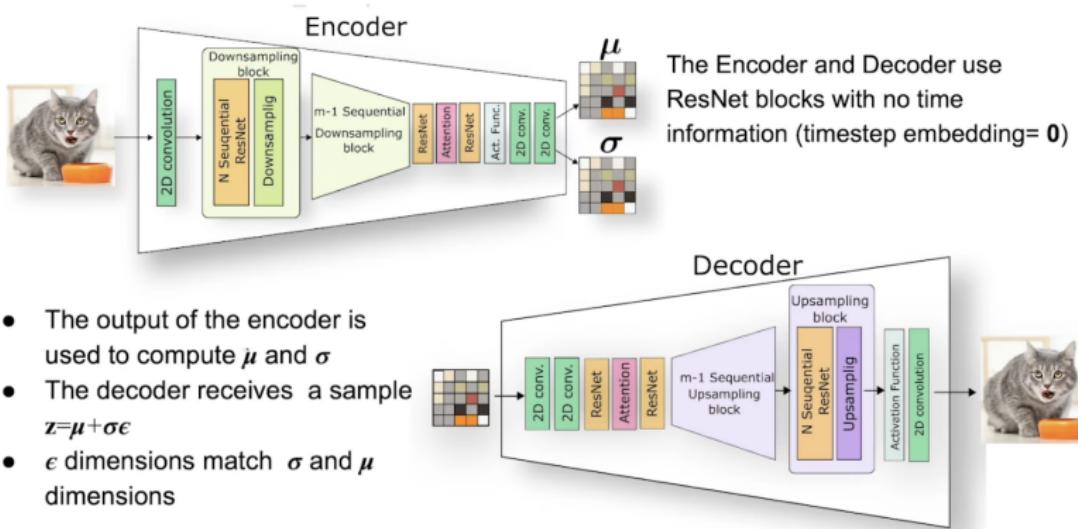


Figure: from Lightning AI video. Note:  $\mu$  and  $\sigma$  can be vector instead of matrices, making implementation easier.

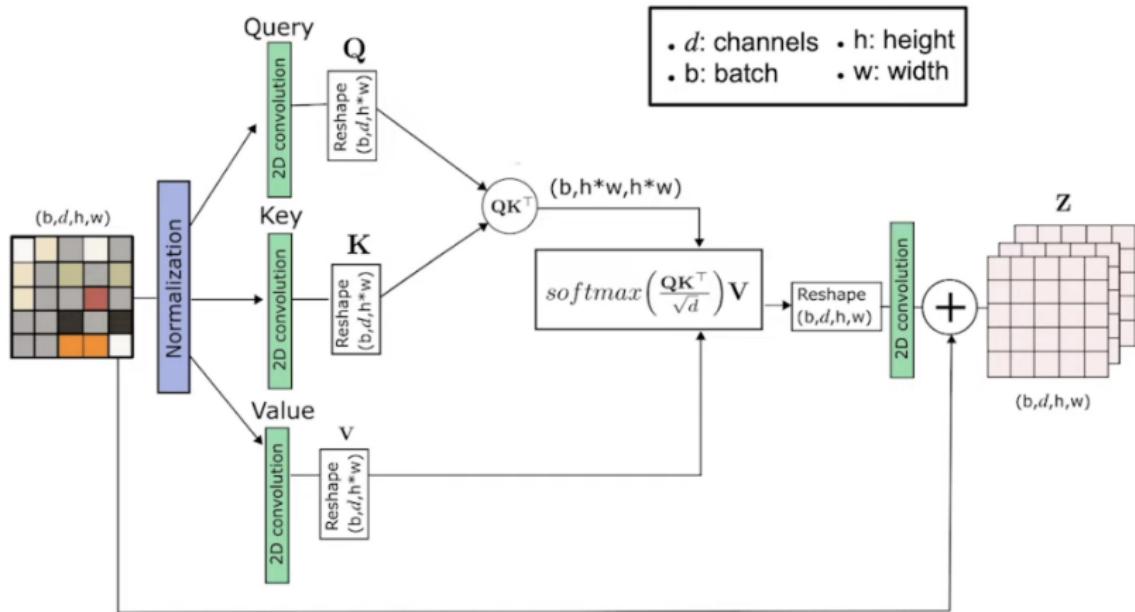


Figure: from Lightning AI video. Attention mechanism in the VAE.

The condition encoder can be arbitrary in theory. e.g. Can be a BERT transformer when condition is text.

The condition can be text or layout, thus they chose to implement the condition encoder as CLIP.

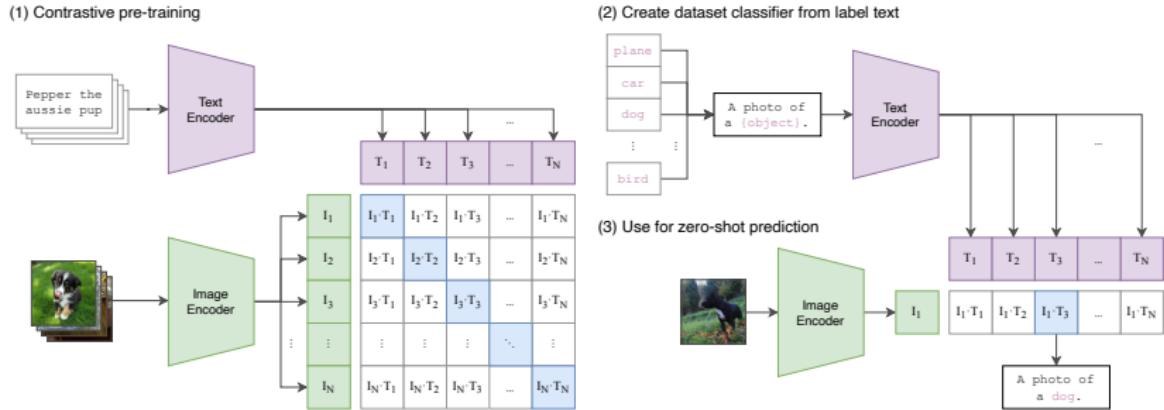


Figure: Standard image models jointly train an image encoder and a linear classifier, whereas CLIP jointly trains an image encoder and a text encoder, to predict the correct pairings of *(image, text)*. Enables zero-shot prediction at inference stage.

Components:

- ▶ Text Encoder: transformers
- ▶ Image Encoder: different variants of ResNets, and Visual Transformer (ViT)
  - ▶ Visual transformers divide images into small blocks processed as “words”.

Besides:

- ▶ Contrastive-learning framework.
- ▶ The model shows in different flavors, also do ensembles to achieve higher performance.
- ▶ Linear probing could be replaced by a classification layer that needs fine-tune.
- ▶ Prompt engineering matters (a lot).

# Quick Recap: CLIP Performance

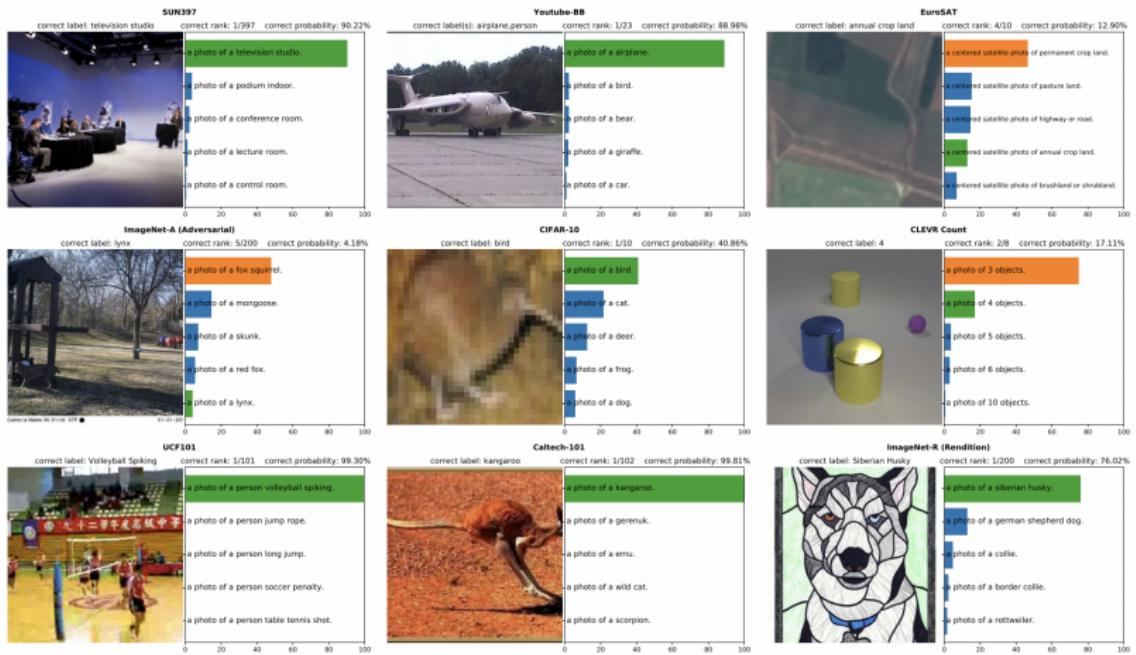


Figure: Zero-shot results. Selected from CLIP paper Figure 21 in their Appendix. Not always good on some fine-grained classification tasks.

- ▶ Scaling: continued improvement on Zero-Shot CLIP to reach overall SOTA performance will require 1000x increase in compute. (current hardware infeasible)
- ▶ Not amazing on all tasks. e.g. trouble with MNIST. Limited by training data set's coverage.
- ▶ No caption generation, only caption retrieval (i.e. select text, not compose new text).
- ▶ Does not address any problem (e.g. data efficiency problem, fairness problem etc.) in deep learning.

- ▶ Image-Searching Engine. (e.g., “car driving in a wood”)
- ▶ Can be used as a discriminator in a GAN framework.

- ▶ Image-Searching Engine. (e.g., “car driving in a wood”)
- ▶ Can be used as a discriminator in a GAN framework.
- ▶ VQ-GAN+CLIP: a powerful model to create image from text. (An Online Tutorial available)
  - ▶ VQ-GAN: generating images that are similar to others (no prompt)
  - ▶ CLIP: text to image
  - ▶ Implementations were made public on Google Colab, no coding needed.
- ▶ etc.

The three components are trained separately. In particular:

1. The autoencoders (VAE  $\mathcal{E}$  and  $\mathcal{D}$ ) are trained first, thus the model knows how to convert images from pixel space to latent space.
2. The denoiser  $\epsilon_\theta$  and condition encoder  $\tau_\theta$  are jointly optimized.

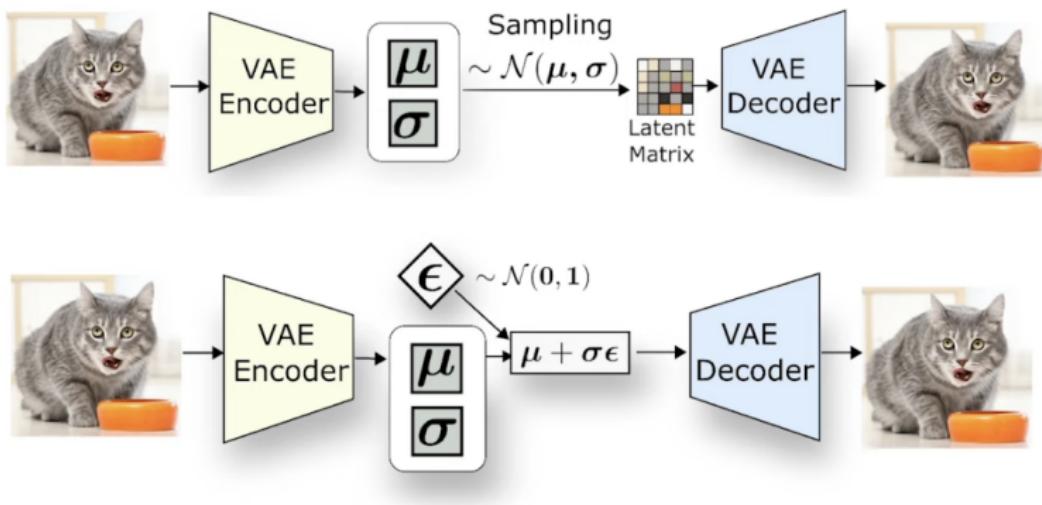


Figure: from Lightning AI video. Re-parameterization trick.

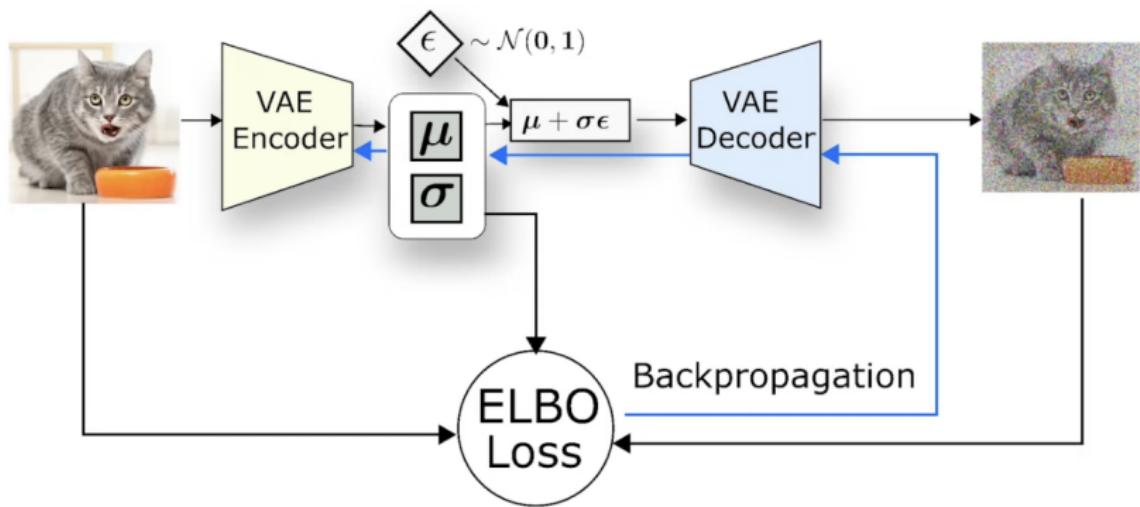


Figure: from Lightning AI video. Evidence Lower Bound Objective (ELBO):  $likelihood(\text{input} - \text{output}) - KLD(\mathcal{N}(\mu, \sigma) || \mathcal{N}(0, 1))$

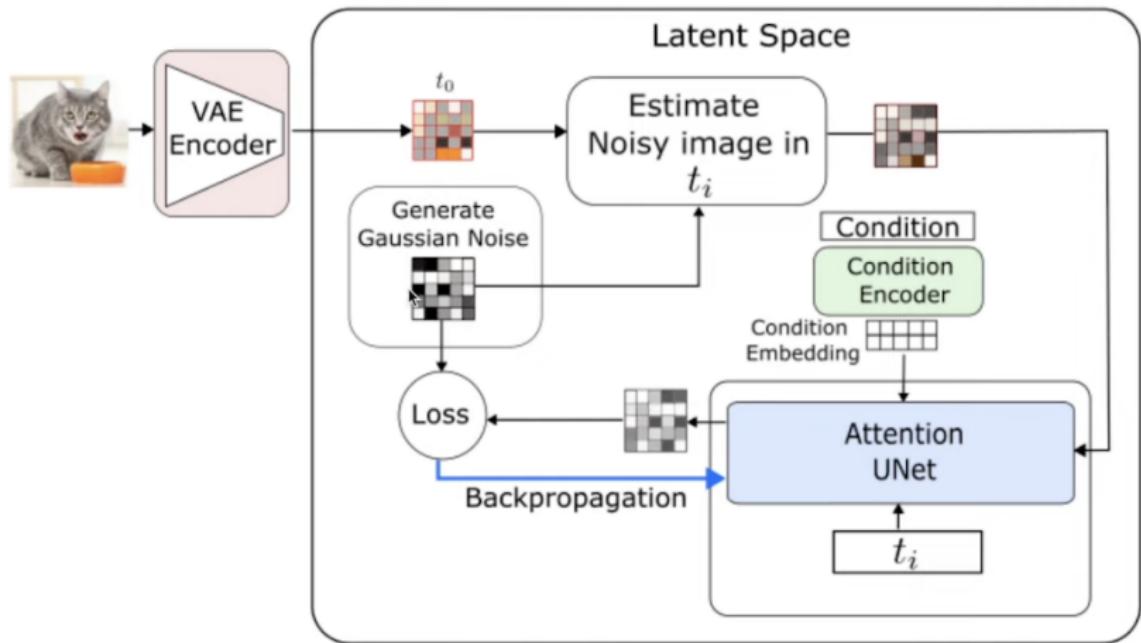


Figure: from Lightning AI video. Loss can be  $\ell_p$  norm.

# Condition Encoder (CLIP) Training

56

Use the cross entropy across images and text to define the loss

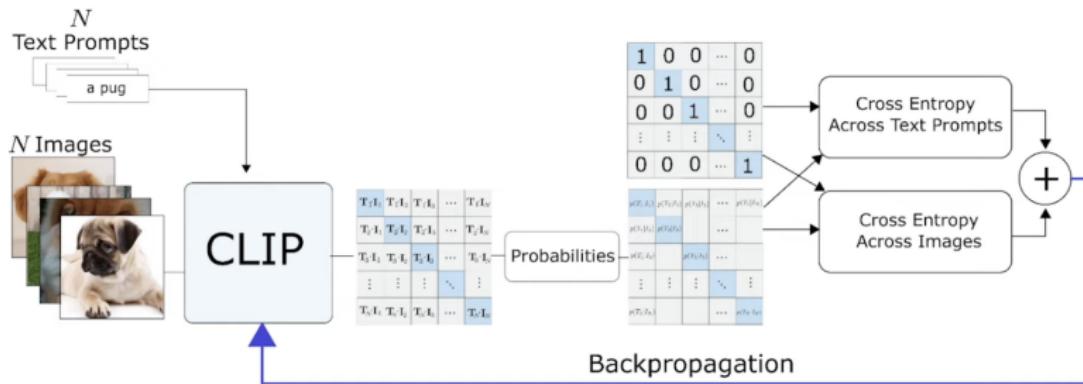


Figure: from Lightning AI video. Sum the loss across rows and columns (i.e., across text and across images) as loss.

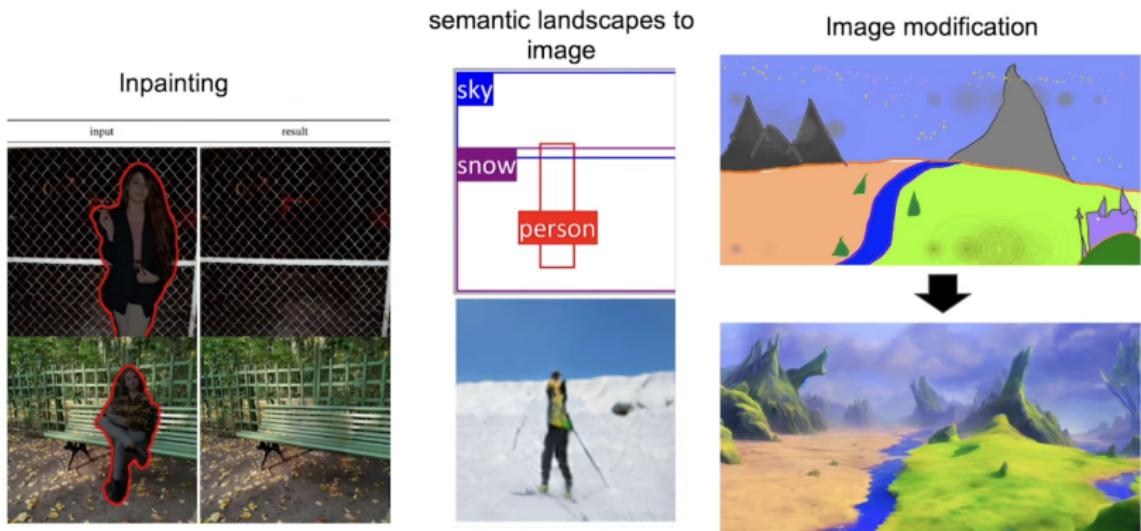


Figure: from Lightning AI video. Achieved via different conditions.  
Inpainting: segmentation map; Semantic Landscape: text explaining the coordinates and objects; Image modification: condition — text, input — original image.



Figure: ImageNet 64 → 256 super-resolution on ImageNet-Val.  
LDM-SR has advantages at rendering realistic textures but SR3 can  
synthesize more coherent fine structures.

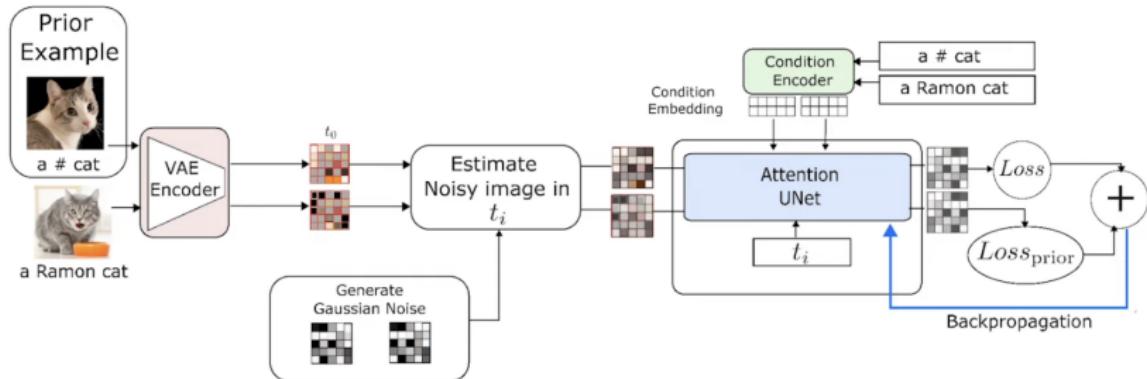


Figure: from Lightning AI video. In this example, a more general sample is provided to avoid language drift. An example of fine-tuning. A few-shot learning scenario.

## Conclusion

---



- ▶ Model follows the Latent Diffusion paper & code (both open-sourced);
- ▶ Teamed up with StabilityAI;
- ▶ Worked with communities e.g. LAION.

For Stable Diffusion:<sup>3</sup>

*The core dataset was trained on LAION-Aesthetics, a soon to be released subset of LAION 5B. LAION-Aesthetics was created with a new CLIP-based model that filtered LAION-5B based on how “**beautiful**” an image was, building on ratings from the alpha testers of Stable Diffusion. LAION-Aesthetics will be released with other subsets in the coming days on <https://laion.ai>.*

Training set matters, a lot.

<sup>3</sup>Official Announcement

DALL-E: <https://arxiv.org/abs/2102.12092>

- ▶ Two stages:
  - ▶ Stage 1: Learning Visual Codebook. Train a discrete variational autoencoder (dVAE) to compress images into tokens.
  - ▶ Stage 2: Learning Prior. Concatenate BPE-encoded text tokens with the image tokens, train an autoregressive transformer to model the joint distribution over the text and image tokens.
  - ▶ Training Procedure: maximizing the evidence lower bound (ELB).
- ▶ gibberish also works: Discovering the Hidden Vocabulary of DALL-E-2

## DALL-E-2

- ▶ Hierarchical Text-Conditional Image Generation with CLIP Latents
- ▶ The three **separated** components to train:
  1. training of CLIP: the most important step;
  2. train the Decoder: generate images based on the CLIP image embedding; Need the trained CLIP from the first step;
  3. train Diffusion Prior Network: takes the CLIP text embeddings and generate the CLIP image embeddings; Need the trained CLIP from the first step.
- ▶ (Unofficial) Code:  
<https://github.com/lucidrains/DALLE2-pytorch>

GLIDE: <https://arxiv.org/abs/2112.10741>

- ▶ Using image down-sampling / up-sampling algorithm to reduce parameter size.
- ▶ Less parameter but slower inference than DALL-E.

GLIDE: <https://arxiv.org/abs/2112.10741>

- ▶ Using image down-sampling / up-sampling algorithm to reduce parameter size.
- ▶ Less parameter but slower inference than DALL-E.

Another Platform mentioned together – MidJourney:

- ▶ An independent research lab  
<https://www.midjourney.com/home/>
- ▶ Users create artwork with Midjourney using Discord bot commands.

Main differences:

- ▶ Used different data sets.
- ▶ LDM works on latent space. GLIDE works on pixel space, reducing parameter size via image down-sampling and up-sampling algorithms.
- ▶ LDM can handle more general conditions, while GLIDE considers only text condition.
- ▶ Different ways of injecting text condition.

Two standard ways of injecting text conditions in DMs:

1. Concatenate the encoded text from a language transformer (usually sentence level, [CLS]) to the image input.
2. Multi-head Cross Attention: Letting **every** U-Net attention layer attend to **all** the text tokens (from language transformer).

GLIDE uses them BOTH, but it is still **not sufficient** in inference time.

- ▶ They used CLIP-guided diffusion to make the generated image better correspond to text.
- ▶ Takes the original de-noised image, and move it towards the direction of “more similar to text goal” (by considering  $\nabla$  of CLIP score).

What is the other way of making the text information more obvious to GLIDE?

At each diffusion step, they applied this trick:

1. Produce the image twice: once with text, once without access to text.
2. Compute the difference between the version with-and-without text, and then move towards the direction of text information.

Isn't it familiar?

Recall LDM:

$$\text{Noise} = \text{Noise}_{uncond} + \beta(\text{Noise}_{cond} - \text{Noise}_{uncond})$$

The trick is the **Classifier-Free Guidance**. (Worked well in practice!)

- ▶ <https://openreview.net/forum?id=qw8AKxfYbI>
- ▶ Classifier Guidance: the diffusion score include the gradient of the log likelihood of an auxiliary classifier model  $p(z|c)$ , where  $c$  is the condition information (e.g. text embedding) and  $z$  is the noisy image. Downside: requires an additional classifier model.
- ▶ Classifier-Free Guidance: jointly train the same architecture with & without the condition, computes a linear combination of the conditional and unconditional score.

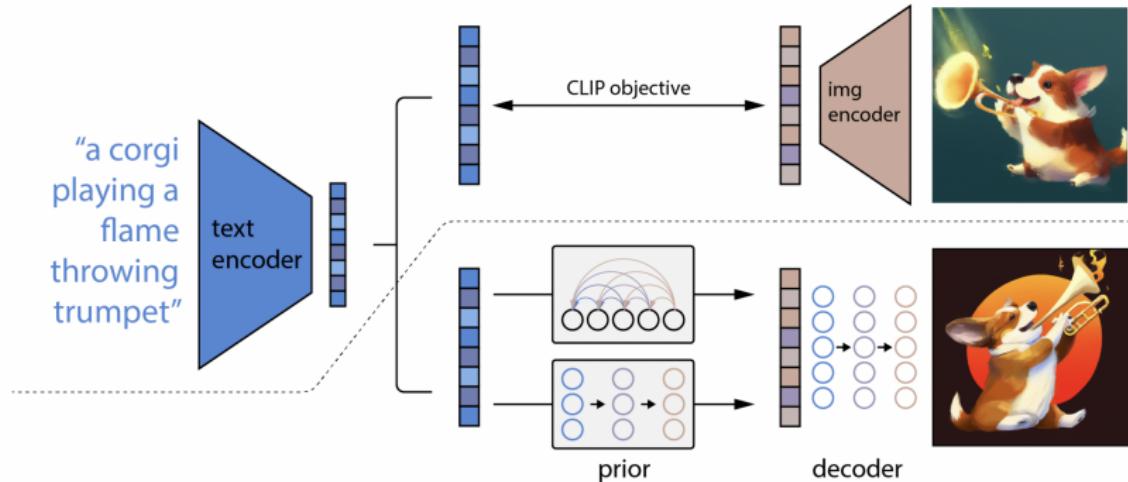


Figure: High-level overview of DALL·E-2. Above the dotted line: training; Below the dotted line: generation.  $z_i$  is the CLIP image embedding,  $y$  is the caption/condition,  $x$  is the image output. Need to train a **prior**  $p(z_i|y)$  (conditioning), a **decoder**  $p(x|z_i, y)$  (implemented as diffusion models).

$$p(x|y) = p(x, z_i|y) = p(x|z_i, y)p(z_i|y)$$

Main differences:

- ▶ Used different data sets.
- ▶ Different architectures:
  - ▶ LDM: VAE, Denoiser, Condition Encoder
  - ▶ DALL·E-2: CLIP, Decoder, Prior (condition encoder)
- ▶ Different ways of injecting condition.

DALL-E-2 Decoder  $p(x|z_i, y)$ , producing image  $x$  from CLIP image embeddings  $z_i$  (and, optionally, text captions  $y$ ).

- ▶ Similar to GLIDE pipeline, but project CLIP embeddings into 4 extra tokens of context, then concatenated to the sequence of outputs from the GLIDE text encoder.  
Classifier-free guidance (with v.s. without CLIP embedding) applied.
- ▶ Also trained two diffusion upsampler models to generate high resolution images. Found it useless to consider  $y$  while upsampling, thus ignored it.

Two versions of DALL·E-2 prior  $p(z_i|y)$ , modeling CLIP image embedding  $z_i$  from caption  $y$ :

- ▶ Autoregressive (AR) prior: <sup>4</sup> converted  $z_i$  into a sequence of discrete codes predicted autoregressively conditioned on  $y$ . Principal Component Analysis (PCA) is applied to reduce the dimensionality of  $z_i$ . Predict the resulting sequence using a Transformer model with a causal attention mask.
- ▶ Diffusion prior: directly model  $z_i$  using a Gaussian diffusion model conditioned on  $y$ . Train a decoder-only Transformer.

$$L_{prior} = \mathbb{E}_{t \sim [1, T], z_i^{(t)} \sim q_t} [\|f_\theta(z_i^{(t)}, t, y) - z_i\|^2]$$

<sup>4</sup>A note on autoregressive v.s. autoencoding.

In addition, both options use CLIP text embedding  $z_t$  since it is a deterministic function of the caption.

- ▶ AR: add  $y$  and  $z_t$  as a prefix to the sequence. Prepend a token  $z_i \cdot z_t$  as well.
- ▶ DM: use  $z_t$  as part of the input sequence of the decoder-only Transformer. Used to improve quality during sampling, generating two samples of  $z_i$  and selecting the one with a higher dot product with  $z_t$ .



Figure: Theatre D'opera Spatial, won the Colorado State Fair digital art contest Aug 2022.

# Social Impact: on the other side

76

PH AI  
@fofrAI

With #stablediffusion img2img, I can help bring my 4yr old's sketches to life.

Baby and daddy ice cream robot monsters having a fun day at the beach. 😊

#AiArtwork



The image shows a side-by-side comparison of two versions of a child's drawing. On the left is the original hand-drawn sketch, which features two characters: one with a large, bulbous head and a grid-like body, and another smaller character with a more organic, spiky appearance. On the right is the AI-generated version, which has transformed these sketches into highly detailed, three-dimensional 3D models. The larger character on the right is now a robot-like figure with a white, segmented head, a metallic body, and articulated arms and legs. The smaller character is also a robot, appearing as a smaller version of the first. They are set against a background of a sandy beach and a bright sky with floating spheres, illustrating how AI can bring a child's imagination to life through digital art.

UCLA

People (who do not have to know AI well) had great fun playing with those tools.

Some artists are greatly upset these days.

- ▶ There has been **serious** ethics concern on the copyright of the online images (to train, we need a lot of images + text descriptions).
- ▶ AI-generated works started to confuse the judge and won prize. [Report from NY Times](#).
- ▶ Anti-AI Movement has been started. See more discussions on [Reddit](#).

Can AI tools replace human artists? What AI art means for human artists? [www.youtube.com/watch?v=sFBfrZ-N3G4](https://www.youtube.com/watch?v=sFBfrZ-N3G4)